

# A Semi Compile/Run-time Map with (Nearly) Zero Overhead Lookup

Fabian Renn-Giles

Watch video of the  
talk here:

[https://www.youtube.com/  
watch?v=gNAbGpV1ZkU](https://www.youtube.com/watch?v=gNAbGpV1ZkU)

[fabian.renn@gmail.com](mailto:fabian.renn@gmail.com)  
@hogliux

# 1. Introduction

# Introduction

## Usage example: C++ ↔ Java Bridge

```
class InputStream // java.io.InputStream
{
public:
    inline void close()          { jclass.getMethod("close").invoke(); }
    inline long skip(long bytes) { return jclass.getMethod("skip").invoke(bytes); }
    inline int read()           { return jclass.getMethod("read").invoke(); }
    .
    .
    .
};
```

called in tight-loop

Ensure getMethod is as fast as possible

Auto-generated C++ code for every Java class

# Introduction

Use a cache!

```
JavaMethod getMethod(const std::string& methodName)
{
    static std::unordered_map<std::string, JavaMethod> cache;
    return cache.try_emplace(methodName, methodName).first->second;
}
```

- try\_emplace needs to calculate the hash of methodName
- Collisions expected as thousands of methods in cache

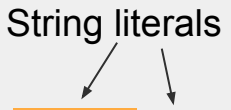
➔ Can we do better?

# Introduction

Use a cache!

```
class InputStream // java.io.InputStream
{
public:
    void close()          { jclass.getMethod("close").invoke(); }
    long skip(long bytes) { return jclass.getMethod("skip").invoke(bytes); }
    int read()            { return jclass.getMethod("read").invoke(); }
};
```

String literals



- methodName is known at compile-time
- it should be possible to calculate the cache slot at compile-time
- Fallback to run-time hash if methodName is not a string literal

➔ **Somebody must have done this already, right?**

# Introduction

Ben Deane's and Jason Turner's “constexpr all the things”<sup>1,2</sup>

```
cx::map<Key, Value>
```

- map is either completely compile-time or completely run-time
- If lookup is compile-time then all values must also be known at compile-time
- JavaMethod can only be resolved at run-time

**✗** Compile-time lookup of run-time value not possible

<sup>1</sup><https://www.youtube.com/watch?v=PJwd4JLYJJY>

<sup>2</sup>there are many more compile-time only maps (for example boost::hana::map)

# Introduction

Louis Dionne's C++ meta-programming talk at meeting C++<sup>1</sup>

```
event_system events{{"foo", "bar", "baz"}}
```

- ✓ Lookup at compile-time and value-storage at runtime
- ✓ Can fallback to run-time lookup if key is not string literal
- ✗ Need to know all possible keys in advance

<sup>1</sup><https://goo.gl/M4vhqc> from slide 73 onwards

# Introduction

Let's roll our own!

```
semi::map<Key, Value>
```

## Requirements

- Lookup at compile-time and value-storage at runtime
- Can fallback to run-time lookup if key is not string literal
- Keys need not to be known ahead of time



## 2. The basic principle

# The basic principle

The basic principle is embarrassingly simple!

```
template <int> std::string map;
```

- ✓ map with `int` keys and `std::string` values
- ✓ Lookup at compile-time and value-storage at runtime
- ✓ Keys need not be known ahead of time

→ **Compile-time lookup, run-time storage map in a single line of code**

# The basic principle

What keys can we use?

```
template <Key> Value map;
```

- Any non-type template parameter types
- essentially `bool`, integral types and `nullptr_t`
- **C++20**: Class types as non-type template parameters!
  - Any compile-time literals, user-defined literals, ...

→ **It's 2018: no compiler supports class types as template parameters**

### 3. Supporting more key types in 2018

# Supporting more key types in 2018

Use typenames instead

```
template <typename> Value map;
```

- We need a “function” which maps any key to a unique typename:

```
UniqueReturnType key2type(Key key);
```

```
std::cout << map<decltype(key2type("foo"))>;
```

➔ Is it possible to write `key2type` in C++?

# Supporting more key types in 2018

## Implementing key2type for integral key types

```
template <auto...> struct dummy_t {};  
  
template <typename Key>  
constexpr auto key2type(Key key)  
{  
    return dummy_t<key>{};  
}
```

```
std::cout << map<decltype(key2type(5))>;
```

- Returns instance of `dummy_t<5>` for key “5”, and so on...

✗ Compiler-error: “key” not a constant expression

# Supporting more key types in 2018

## Implementing key2type for integral key types

```
template <auto...> struct dummy_t {};  
  
template <typename Lambda>1  
constexpr auto key2type(Lambda lambda)  
{  
    return dummy_t<lambda()>{};  
}
```

```
std::cout << map<decltype(key2type([]() { return 5; })))>;2
```

<sup>1</sup>enable\_if template parameter removed for brevity

<sup>2</sup>only shown to demonstrate usage: error due to lambda being used in an unevaluated context - will be resolved on slide 22

# Supporting more key types in 2018

## Implementing key2type for integral key types

```
template <auto...> struct dummy_t {};  
  
template <typename Lambda>1  
constexpr auto key2type(Lambda lambda)  
{  
    return dummy_t<Lambda()>{};  
}  
#define ID(x) []() constexpr { return x; }
```

```
std::cout << map<decltype(key2type(ID(5)))> << std::endl;2
```

- Still doesn't work with string literals etc.

<sup>1</sup>enable\_if template parameter removed for brevity

<sup>2</sup>only shown to demonstrate usage: error due to lambda being used in an unevaluated context - will be resolved on slide 22



# Supporting more key types in 2018

## Supporting string literals:

```
template <typename Lambda, std::size_t... I>
constexpr auto str2type(Lambda lambda, std::index_sequence<I...>)
{
    return dummy_t<lambda()[I]...>{};
}

template <typename Lambda>1
constexpr auto key2type(Lambda lambda)
{
    return array2type (lambda, std::make_index_sequence<strlen3(lambda())>{});
}
```

```
std::cout << map<decltype(key2type(ID("foo")))> << std::endl;2
```

- Maps “foo” onto dummy\_t<'f', 'o', 'o'>

<sup>1</sup>enable\_if template parameter removed for brevity

<sup>2</sup>only shown to demonstrate usage: error due to lambda being used in an unevaluated context - will be resolved on slide 22

<sup>3</sup>strlen is marked constexpr in gcc only: substitute implementation for other compilers is left as an exercise for the reader

## 4. Putting it all together

# Putting it all together

```
template <typename> std::string map_value;
```

```
map_value<decltype(key2type(ID("conference")))> = "cppcon";
```

# Putting it all together

```
template <typename Value typename>  
Value& static_map_get()  
{  
    static Value value;  
    return value;  
}
```

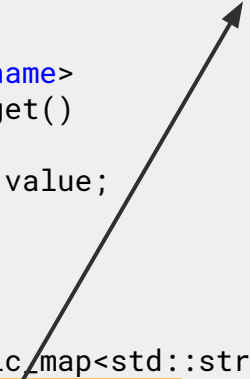
```
static_map_get<std::string, decltype(key2type(ID("conference")))>() = "cppcon";
```

# Putting it all together

```
template <typename Key, typename Value>
class static_map
{
```

```
public:
    template <typename>
    static Value& get()
    {
        static Value value;
        return value;
    }
};
```

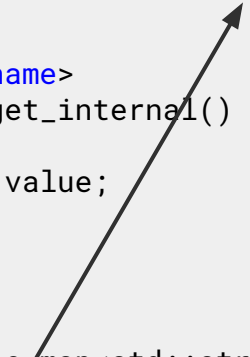
```
using map = static_map<std::string, std::string>;
map::get<decltype(key2type(ID("conference")))>() = "cppcon";
```



# Putting it all together

```
template <typename Key, typename Value>
class static_map
{
public:
    template <typename Lambda>
    static Value& get(Lambda lambda)
    {
        return get_internal<decltype(key2type(lambda))>();
    }
private:
    template <typename>
    static Value& get_internal()
    {
        static Value value;
        return value;
    }
};

using map = static_map<std::string, std::string>;
map::get(ID("conference")) = "cppcon";1
```



<sup>1</sup>constexpr lambda is now being evaluated as a normal parameter which resolves the issue mentioned in the footnote of slides 15,16,17

# Putting it all together

```
template <typename Key, typename Value>
class static_map
{
public:
    template <typename Lambda>
    static Value& get(Lambda lambda)
    {
        static_assert(std::is_convertible_v<decltype(lambda()), Key>);
        return get_internal<decltype(key2type(lambda))>();
    }
private:
    template <typename>
    static Value& get_internal()
    {
        static Value value;
        return value;
    }
};

using map = static_map<std::string, std::string>;
map::get(ID("conference")) = "cppcon";1
```

<sup>1</sup>constexpr lambda is now being evaluated as a normal parameter which resolves the issue mentioned in the footnote of slides 15,16,17

# Putting it all together

What's the performance?

```
int& getAge()
{
    using map =
        semi::static_map<std::string, int>;

    return map::get(ID("age"));
}
```

```
mov eax, OFFSET FLAT:dummy_tlvalue
ret
```

Compiler Explorer: <https://gcc.godbolt.org/z/6CRZ9W>

➔ Just as fast as accessing a global variable



## 5. Optional run-time lookup

# Optional run-time lookup

Impossible to write run-time version of key2type...

```
UniqueReturnType key2type(Key key);
```

...also impossible to somehow call get with run-time key

```
template <typename Lambda> get(Lambda lambda);
```

We can write a big switch-like statement:

```
if (key == "food") return get(ID("food"));  
else if (key == "drink") return get(ID("drink"));  
...
```

...but then all keys need to be known in advance.

# Optional run-time lookup


Solution: maintain run-time map of pointers to static locals

```
std::unordered_map<Key, Value*> runtime_map;
```

Add to map when value is accessed via compile-time literal:

```
template <typename>
static Value& get_internal (const Key& key)
{
    static Value value;
    runtime_map[key] = &value;
    return value;
}
```

**Slow:** Executed  
every time get is  
called



➔ Only add to map when `get_internal` is executed for the first time?

# Optional run-time lookup

Closer look at the static local variable:


```
template <typename>
static Value& get_internal()
{
    static Value value;

    return value;
}
```

# Optional run-time lookup

Closer look at the static local variable:

```
template <typename>
static Value& get_internal()
{
    static Value value;
    return value;
}
```



- 1. Creates storage
- 2. Calls constructor

# Optional run-time lookup

Closer look at the static local variable:

```
struct ConstructorInvoker
{
    ConstructorInvoker(char* mem) { new (mem) Value; }
};

template <typename>
static Value& get_internal()
{
    alignas (Value) static char storage[sizeof(Value)];
    static ConstructorInvoker invoker(storage);

    return *reinterpret_cast<Value*> (storage);
}
```

➔ This change has no effect on the generated assembly

# Optional run-time lookup

Undefined behaviour?!?

```
struct ConstructorInvoker
{
    ConstructorInvoker(char* mem) { new
};

template <typename>
static Value& get_internal()
{
    alignas (Value) static char storage[sizeof(Value)];
    static ConstructorInvoker invoker(storage);

    return *reinterpret_cast<Value*> (storage);
}
```



- No alignment issues
  - No aliasing
- **No excuses: standard says `reinterpret_cast` to `Value` is UB!**

# Optional run-time lookup

Undefined behaviour?!?

```
struct ConstructorInvoker
{
    ConstructorInvoker(char* mem) { new (mem) Value; }
};

template <typename>
static Value& get_internal()
{
    alignas (Value) static char storage[sizeof(Value)];
    static ConstructorInvoker invoker(storage);

    return *std::launder(reinterpret_cast<Value*> (storage));
}
```

Meet the most obscure C++ function: `std::launder`

- Can be used to avoid UB in these types of cases
- [cppreference.com](http://cppreference.com) shows exactly this example



# Optional run-time lookup

Controlling the construction with an init flag

```
struct ConstructorInvoker
{
    ConstructorInvoker(char* mem) { new (mem) Value; }
};

template <typename>
static Value& get_internal(const Key& key)
{
    alignas (Value) static char storage[sizeof(Value)];
    static ConstructorInvoker invoker(storage);

    return *reinterpret_cast<Value*> (storage);
}
```

# Optional run-time lookup

Controlling the construction with a needs\_init flag

```
void initialise(const Key& key, char* mem, bool& needs_init);

template <typename>
static Value& get_internal(const Key& key)
{
    alignas (Value) static char storage[sizeof(Value)];
    static bool needs_init = true;

    if (needs_init) {
        initialise(key, storage, needs_init); needs_init = false;
    }

    return *reinterpret_cast<Value*> (storage);
}
```

No more thread-safety

- Generated code is faster: no more locks
- Just like STL containers, we do not claim that `semi::map` is thread-safe

# Optional run-time lookup

Controlling the construction with a needs\_init flag

```
void initialise(const Key& key, char* mem, bool& needs_init);

template <typename>
static Value& get_internal(const Key& key)
{
    alignas (Value) static char storage[sizeof(Value)];
    static bool needs_init = true;

    if (__builtin_expect(needs_init, false)) {
        initialise(key, storage, needs_init); needs_init = false;
    }

    return *reinterpret_cast<Value*> (storage);
}
```

No more thread-safety

- Generated code is faster: no more locks
- Just like STL containers, we do not claim that `semi::map` is thread-safe

# Optional run-time lookup

## Closer look at the `initialise` function

```
void initialise(const Key& key, char* mem, bool& needs_init)
{
    new (mem) Value;
}
```

# Optional run-time lookup

## Closer look at the `initialise` function

```
void initialise(const Key& key, char* mem, bool& needs_init)
{
    runtime_map.try_emplace(key, new (mem) Value
                                );
}

static std::unordered_map<Key, std::unique_ptr<Value> > runtime_map;
```

- Add `Value` to `runtime_map`

# Optional run-time lookup

## Closer look at the initialise function

```
void initialise(const Key& key, char* mem, bool& needs_init)
{
    runtime_map.try_emplace(new (mem) Value, ValueDeleter{needs_init});
}

struct ValueDeleter
{
    void operator()(Value* v) { v->~Value(); }
    bool& needs_init;
};

static std::unordered_map<Key, std::unique_ptr<Value, ValueDeleter>> runtime_map;
```

- Add Value to runtime\_map
- Ensure that destructor is called but that storage is not deallocated

# Optional run-time lookup

## Closer look at the `initialise` function

```
void initialise(const Key& key, char* mem, bool& needs_init)
{
    runtime_map.try_emplace(new (mem) Value, ValueDeleter{needs_init});
}

struct ValueDeleter
{
    void operator()(Value* v) { v->~Value(); needs_init = true; }
    bool& needs_init;
};

static std::unordered_map<Key, std::unique_ptr<Value, ValueDeleter>> runtime_map;
```

- Add `Value` to `runtime_map`
- Ensure that destructor is called but that storage is not deallocated

# Optional run-time lookup

Getting a value from a runtime key is now trivial!

```
static Value& get(const Key& key)
{
    return *runtime_map[key];
}

using map = semi::static_map<std::string, std::string>

std::string key;
std::cin >> key;
map::get(key);
```

- We also need to account for the case where a value is accessed via run-time key first and **then** via a compile-time key literal
- Not very interesting: see full-code for details on how this is done



# Optional run-time lookup

Additional benefit: we can now also remove values from our map at run-time

```
static void clear()
{
    runtime_map.clear();
}

template <typename Lambda>
static void erase(Lambda lambda)
{
    runtime_map.erase(lambda());
}

static std::unordered_map<Key, std::unique_ptr<Value, ValueDeleter>> runtime_map;
```

➔ Removes both from runtime\_map and compile-time map

# Optional run-time lookup

And what about performance?

```
int& getAge()
{
    using map =
        semi::static_map<std::string, int>;

    return map::get(ID("age"));
}
```

```
cmpb $0, map::needs_init
jne initialise
movl map::storage, %eax
ret
```

`.initialise:`

`. Lots of code`

`.`

} Only  
executed the  
first time  
value is  
accessed

Compiler Explorer: <https://gcc.godbolt.org/z/g7Rb4Z>

6. But everything is static?

# But everything is static?

```
template <typename Key, typename Value>
class static_map
{
public:
    static_map() = delete;

    template <typename Lambda>
    static Value& get(Lambda lambda);

    template <typename Lambda>
    static void erase(Lambda lambda);

    static void clear();
    ...
};
```

- Storage needs to be static!
- So all methods declared static as well
- No per-instance storage

# But everything is static?

Do you really need a non-static map?

- Caches are often used in singletons anyway

```
JavaMethod getMethod(const std::string& methodName)
{
    using cache = semi::static_map<std::string, JavaMethod> cache;
    return cache::get(methodName);
}
```

Collision if semi-map with same Key/Value types exists anywhere else in the code

- Life-time can still be managed with clear method

# But everything is static?

Do you really need a non-static map?

- Caches are often used in singletons anyway

```
JavaMethod getMethod(const std::string& methodName)
{
    struct Tag {};
    using cache = semi::static_map<std::string, JavaMethod, Tag> cache;


    return cache::get(methodName);
}
```

- Life-time can still be managed with clear method

➔ But I still really need a non-static map - is this possible?

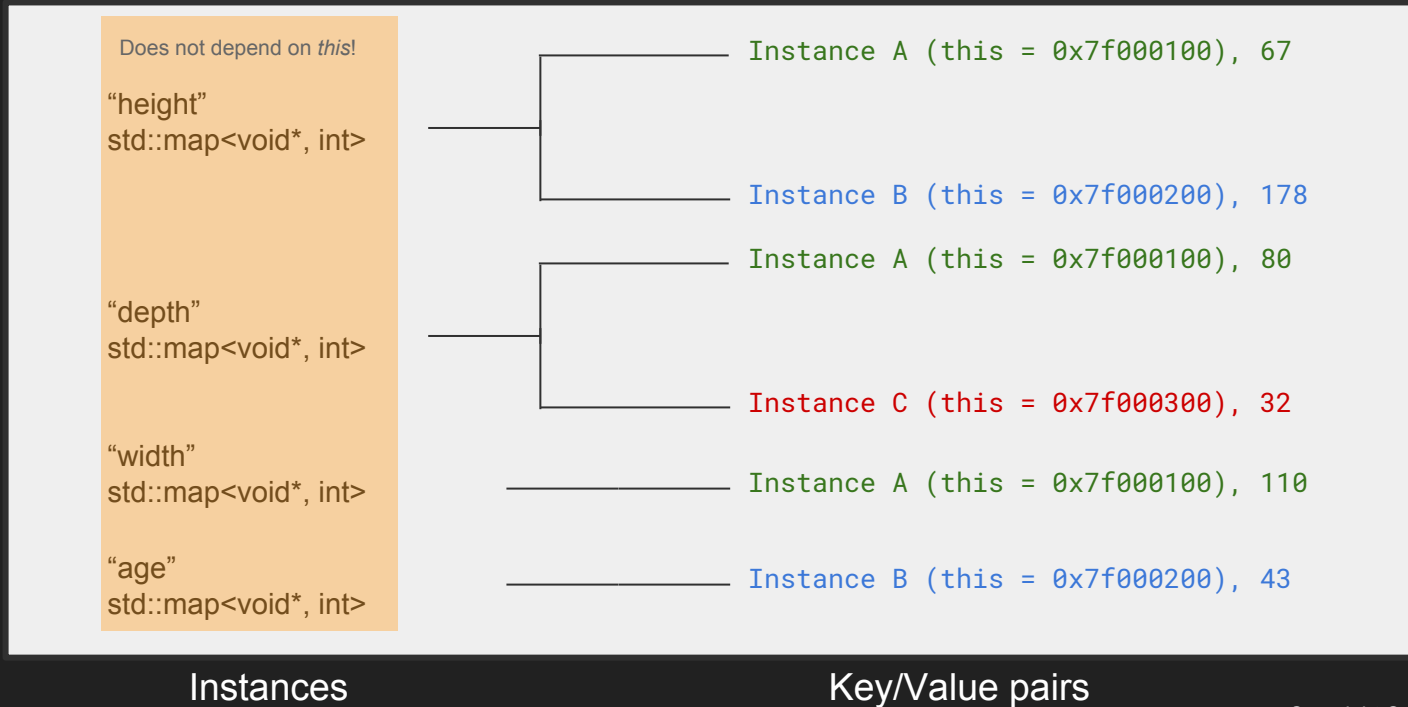
# But everything is static?

Normal runtime-maps: **each this pointer points to a map of keys**



# But everything is static?

Invert the map: **each key points to a map of this pointers!**





# But everything is static?

Solution: **each key points to a map of this pointers!**

```
using map = semi::static_map<Key, Value>;
```

# But everything is static?

**Solution: each key points to a map of this pointers!**

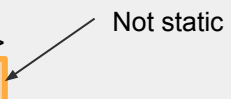
```
using map = semi::static_map<Key, std::map<void*, Value>>;
```

# But everything is static?

Solution: each key points to a map of this pointers!

```
namespace semi
{
    template <typename Key, typename Value>
    class map
    {
    public:
        template <typename Lambda>
        Value& get(Lambda lambda)
        {
            return instance_map::get(lambda)[this];
        }
    private:
        using instance_map = semi::static_map<Key, flat_map<void*, Value>>;
    };
} // namespace semi

semi::map<std::string, std::string> m;
m.get(ID("conference")) = "cppcon";
```



# But everything is static?

Invert the map: each key points to a map of this pointers!

```
namespace semi
{
template <typename Key, typename Value>
class map
{
public:
    template <typename Lambda>
    Value& get(Lambda lambda)
    {
        return instance_map::get(lambda)[this];
    }
private:
    using instance_map = semi::static_map<Key, std::map<void*, Value>>;
};
} // namespace semi

semi::map<std::string, std::string> m;
m.get(ID("conference")) = "cppcon";
```

Uses static\_map


# But everything is static?

Invert the map: each key points to a map of this pointers!

```
namespace semi
{
template <typename Key, typename Value>
class map
{
public:
    template <typename Lambda>
    Value& get(Lambda lambda)
    {
        return instance_map::get(lambda, [this]);
    }
private:
    using instance_map = semi::static_map<Key, std::map<void*, Value>>;
};
} // namespace semi

semi::map<std::string, std::string> m;
m.get(ID("conference")) = "cppcon";
```

Use this parameter as the key




# But everything is static?

But wait, doesn't that defeat the object?

- Calculating the hash of a memory address is often computationally more efficient compared to hashing other key types
- You normally have a large number of keys compared to the number of instances (especially true for caches)

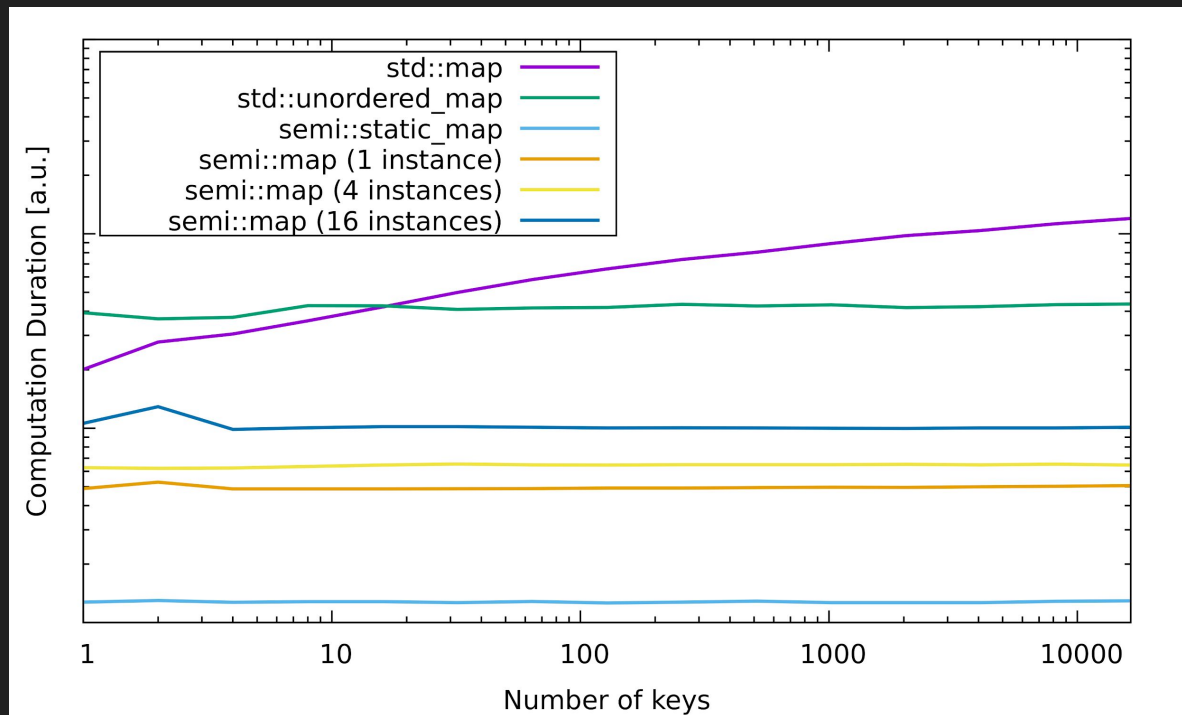
```
using map = semi::static_map<Key, std::map<void*, Value>1>;
```

Often only has 1-3 entries



# But everything is static?


## Benchmarks



Using: `clang++ -Ofast -march=native`

# But everything is static?

## Benchmarks





# Thank you!



Code: <https://github.com/hogliux/semimap>


Video: <https://www.youtube.com/watch?v=qNAbGpV1ZkU>

@hogliux  
[fabian.renn@gmail.com](mailto:fabian.renn@gmail.com)

# Optional run-time lookup

Only add Value pointer the first time get\_internal is called

```
struct ValueWrapper
{
    ValueWrapper (const Key& key) { runtime_map[key] = &v; }
    Value v;
};
```

**Faster:** Only executed the first time get\_internal is called! 

```
template <typename>
static Value& get_internal (const Key& key)
{
    static ValueWrapper value(key);

    return value.v;
}
```

# Optional run-time lookup

Getting a value from a runtime key is now trivial!

```
static Value& get(const Key& key)
{
    auto it = runtime_map.find(key);
    if (it != runtime_map.end())
        return it->second;
    return runtime_map.emplace_hint(it, key, {new Value ValueDeleter{nullptr}})->second;
}
```

Same as before!

- Allocate value on heap if not already in runtime\_map
- We don't have an `init_flat` so pass `nullptr`

# Managing key/value life-time

## Closer look at the initialise function

```
void initialise(const Key& key, char* mem, bool& needs_init)
{
    auto it = runtime_map.find(key);

    if (it == runtime_map.end())
        runtime_map.try_emplace(key, new (mem) Value, ValueDeleter{&needs_init});
    else
        it->second
            = std::unique_ptr<Value, ValueDeleter> (new (mem) Value (std::move(it->second)),
                                                    ValueDeleter{&needs_init});
}

struct ValueDeleter
{
    void operator()(Value* v)
    { if (needs_init != nullptr) { v->~Value(); *needs_init = true;
      else { delete v; } }
    }
    bool* needs_init = nullptr;
};
```

# Putting it all together

What's the performance?

```
std::string& getName()  
{  
    using map =  
        semi::static_map<std::string, std::string>;  
    return map::get(ID("name"));  
}
```

```
movzx eax, BYTE PTR guard variable  
test al, al  
je .L15  
mov eax, OFFSET FLAT:dummy_tlvalue  
ret  
.L15:  
sub rsp, 8  
mov edi, OFFSET FLAT:guard variable  
call __cxa_guard_acquire  
.  
.  
.
```

Only  
executed the  
first time  
value is  
accessed

Compiler Explorer: <https://gcc.godbolt.org/z/Kl8Ynk>

➔ Like any static local variable with non-default constructor,  
the compiler will generate lock guards