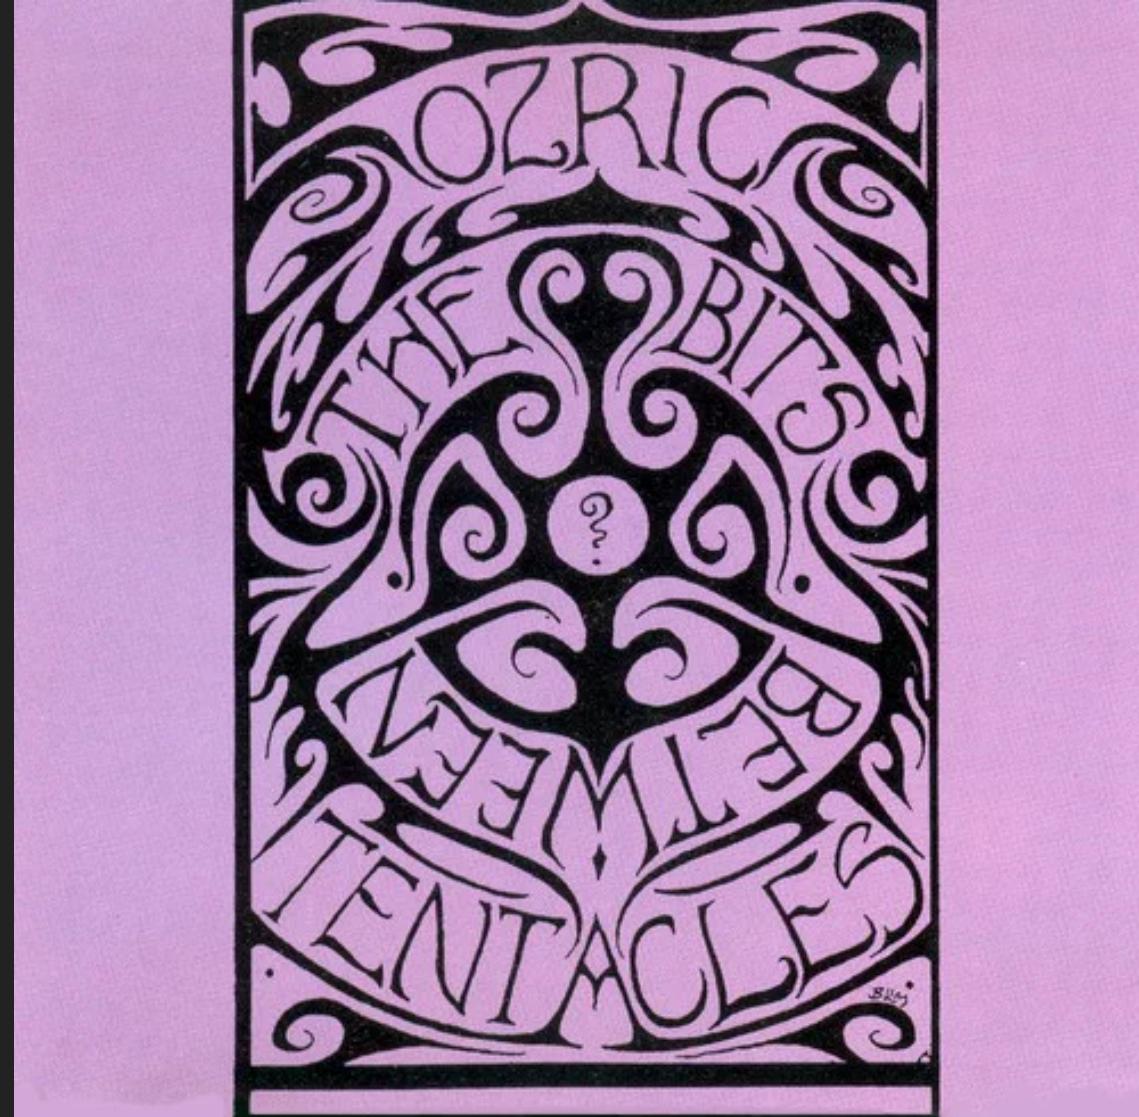


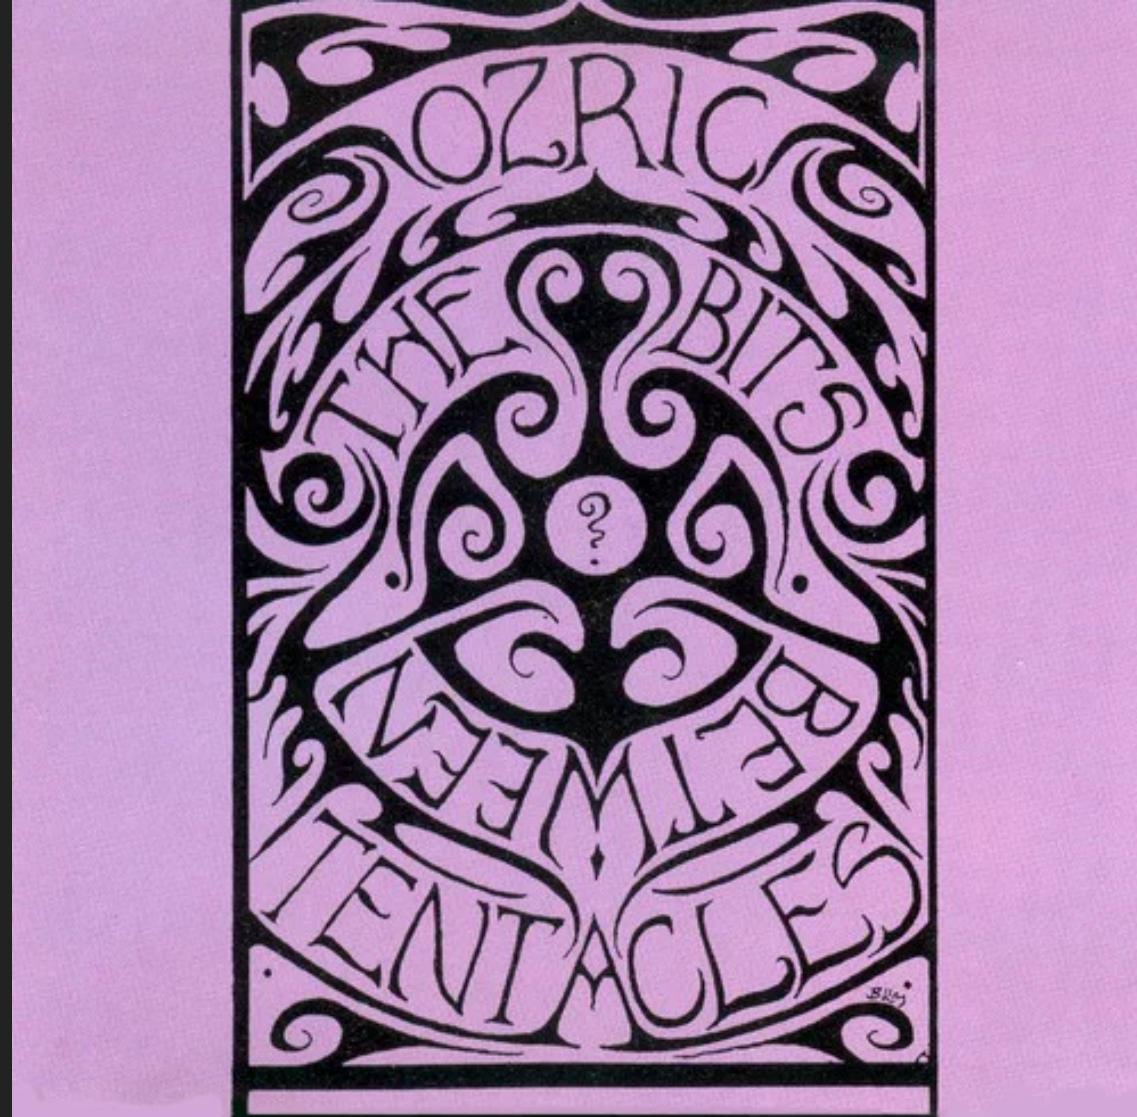
# THE BITS BETWEEN THE BITS

## HOW WE GET TO `main()`

MATT GODBOLT  
@MATTGODBOLT

CPPCON 2018





*"Good, but non-essential"* — [progarchives.com](http://progarchives.com)

# A PROGRAM

```
int main() {}
```

# A PROGRAM

```
int main() {}
```

```
$ gcc -Os empty.c -o c/empty
$ g++ -Os empty.cpp -o cpp/empty
```

# A PROGRAM

```
int main() {}  
  
$ gcc -Os empty.c -o c/empty  
$ g++ -Os empty.cpp -o cpp/empty  
  
$ ls -l c/empty cpp/empty
```

# A PROGRAM

```
int main() {}
```

```
$ gcc -Os empty.c -o c/empty
$ g++ -Os empty.cpp -o cpp/empty

$ ls -l c/empty cpp/empty
7976 c/empty
```

# A PROGRAM

```
int main() {}
```

```
$ gcc -Os empty.c -o c/empty
$ g++ -Os empty.cpp -o cpp/empty
```

```
$ ls -l c/empty cpp/empty
7976 c/empty
7976 cpp/empty
```

# WHAT'S IN IT?

# WHAT'S IN IT?

```
$ objdump --no-show-raw-instr -dC cpp/empty
```

# WHAT'S IN IT?

```
$ objdump --no-show-raw-instr -dC cpp/empty
```

```
research/out/empty/c++/none/dynamic/empty:      file format elf64-x86-64
```

```
Disassembly of section .init:
```

```
00000000004003b0 <_init>:  
 4003b0: sub    $0x8,%rsp  
 4003b4: mov    0x200c3d(%rip),%rax          # 600ff8 <__gmon_start__>  
 4003bb: test   %rax,%rax  
 4003be: je     4003c2 <_init+0x12>  
 4003c0: callq  *%rax  
 4003c2: add    $0x8,%rsp  
 4003c6: retq
```

```
Disassembly of section .text:
```

# WHAT'S IN IT?

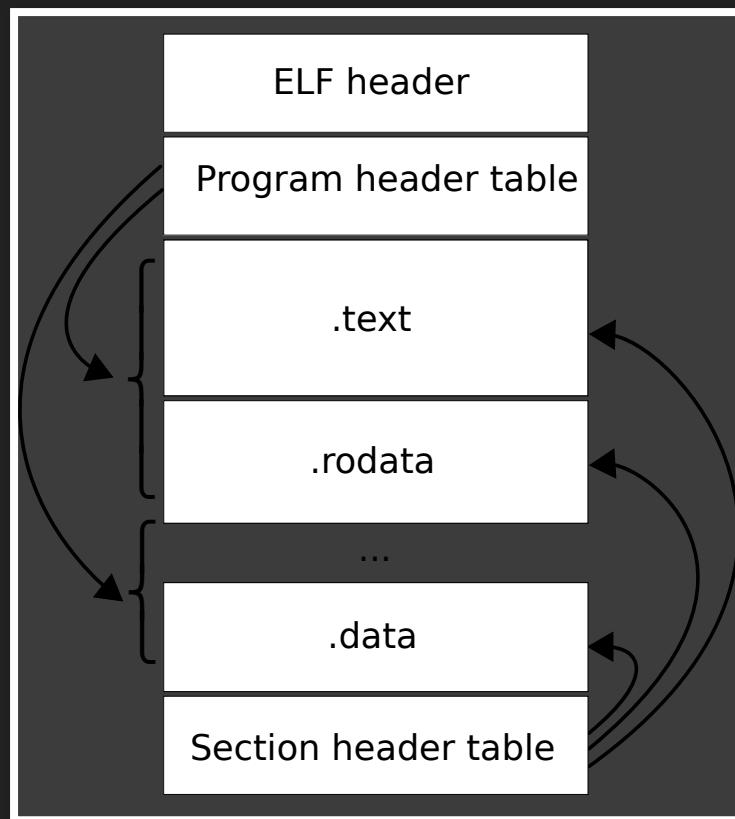
```
$ readelf -a cpp/empty
```

# WHAT'S IN IT?

```
$ readelf -a cpp/empty
```

```
ELF Header:  
  Magic: 7f 45 4c 46 02 01 01 00 00 00 00 00 00 00 00 00  
  Class: ELF64  
  Data: 2's complement, little endian  
  Version: 1 (current)  
  OS/ABI: UNIX - System V  
  ABI Version: 0  
  Type: EXEC (Executable file)  
  Machine: Advanced Micro Devices X86-64  
  Version: 0x1  
  Entry point address: 0x4003e0  
  Start of program headers: 64 (bytes into file)  
  Start of section headers: 6248 (bytes into file)  
  Flags: 0x0  
  Size of this header: 64 (bytes)
```

# THE ELF FILE FORMAT



By Surueña [CC BY-SA 3.0], from Wikimedia Commons

# SECTIONS

- `.text` – code
- `.rodata` – read-only data
- `.data` – read/write data
- `.bss` – zero-initialised data

# HOW WE GET TO `main()`

A (slightly) more interesting program

```
struct Foo {  
    static int numFoos;  
    Foo() {  
        numFoos++;  
    }  
    ~Foo() {  
        numFoos--;  
    }  
};  
int Foo::numFoos;
```

```
Foo globalFoo;  
  
int main() {  
    std::cout << "numFoos = "  
        << Foo::numFoos << "\n";  
}
```

# WHAT DOES IT PRINT?

```
$ g++ -O0 -g global.cpp -o global  
$ ./global
```

# WHAT DOES IT PRINT?

```
$ g++ -O0 -g global.cpp -o global
$ ./global
numFoos = 1
```

# CODE ARCHAEOLOGY - PART 1



# CALL STACK

```
#0 Foo::Foo (this=0x601050 <global>) at global.cpp:6
#1 0x000000000040079d in __static_initialization_and_destruction_0 (
    __initialize_p=1, __priority=65535) at global.cpp:14
#2 0x00000000004007b3 in __GLOBAL__sub_I_global.cpp(void) () at global.cpp:18
#3 0x000000000040082d in __libc_csu_init ()
#4 0x00007ffff70c1b28 in __libc_start_main (main=0x400702 <main()>, argc=1,
    ... at ../csu/libc-start.c:266
#5 0x000000000040064a in __start ()
```

# CALL STACK

```
#0 Foo::Foo (this=0x601050 <global>) at global.cpp:6
#1 0x000000000040079d in __static_initialization_and_destruction_0 (
    __initialize_p=1, __priority=65535) at global.cpp:14
#2 0x00000000004007b3 in __GLOBAL__sub_I_global.cpp(void) () at global.cpp:18
#3 0x000000000040082d in __libc_csu_init ()
#4 0x00007ffff70c1b28 in __libc_start_main (main=0x400702 <main()>, argc=1,
    ... at ../csu/libc-start.c:266
#5 0x000000000040064a in __start ()
```

# CALL STACK

```
#0 Foo::Foo (this=0x601050 <global>) at global.cpp:6
#1 0x000000000040079d in __static_initialization_and_destruction_0 (
    __initialize_p=1, __priority=65535) at global.cpp:14
#2 0x00000000004007b3 in __GLOBAL__sub_I_global.cpp(void) () at global.cpp:18
#3 0x000000000040082d in __libc_csu_init ()
#4 0x00007ffff70c1b28 in __libc_start_main (main=0x400702 <main()>, argc=1,
    ... at ../csu/libc-start.c:266
#5 0x000000000040064a in __start ()
```

# CALL STACK

```
#0  Foo::Foo (this=0x601050 <global>) at global.cpp:6
#1  0x000000000040079d in __static_initialization_and_destruction_0 (
    __initialize_p=1, __priority=65535) at global.cpp:14
#2  0x00000000004007b3 in __GLOBAL__sub_I_global.cpp(void) () at global.cpp:18
#3  0x000000000040082d in __libc_csu_init ()
#4  0x00007ffff70c1b28 in __libc_start_main (main=0x400702 <main()>, argc=1,
    ... at ../csu/libc-start.c:266
#5  0x000000000040064a in _start ()
```

# CALL STACK

```
#0  Foo::Foo (this=0x601050 <global>) at global.cpp:6
#1  0x000000000040079d in __static_initialization_and_destruction_0 (
    __initialize_p=1, __priority=65535) at global.cpp:14
#2  0x00000000004007b3 in __GLOBAL__sub_I_global.cpp(void) () at global.cpp:18
#3  0x000000000040082d in __libc_csu_init ()
#4  0x00007ffff70c1b28 in __libc_start_main (main=0x400702 <main()>, argc=1,
    ... at ../csu/libc-start.c:266
#5  0x000000000040064a in __start ()
```

# WHERE DO THESE FUNCTIONS COME FROM?

# WHERE DO THESE FUNCTIONS COME FROM?

[https://godbolt.org/z/cppcon2018\\_1](https://godbolt.org/z/cppcon2018_1)

# WHO CALLS THIS FUNCTION?

# WHO CALLS THIS FUNCTION?



*By Google Inc. - [https://chromium.googlesource.com/chromium/src/+master/ui/webui/resources/images/google\\_logo.svg](https://chromium.googlesource.com/chromium/src/+master/ui/webui/resources/images/google_logo.svg), Public Domain, [Link](#)*

# LIBC SPELUNKING



# LIBC SPELUNKING

```
// Paraphrased from glibc/csu/elf-init.c
typedef void (*init_func)(int, char **, char **);

extern init_func __init_array_start[];
extern init_func __init_array_end[];

int __libc_csu_init(int argc, char **argv, char **envp) {
    const size_t size = __init_array_end - __init_array_start;
    for (size_t i = 0; i < size; i++)
        (*__init_array_start[i])(argc, argv, envp);
}
```

# LIBC SPELUNKING

```
// Paraphrased from glibc/csu/elf-init.c
typedef void (*init_func)(int, char **, char **);

extern init_func __init_array_start[];
extern init_func __init_array_end[];

int __libc_csu_init(int argc, char **argv, char **envp) {
    const size_t size = __init_array_end - __init_array_start;
    for (size_t i = 0; i < size; i++)
        (*__init_array_start[i])(argc, argv, envp);
}
```

# LIBC SPELUNKING

```
// Paraphrased from glibc/csu/elf-init.c
typedef void (*init_func)(int, char **, char **);

extern init_func __init_array_start[];
extern init_func __init_array_end[];

int __libc_csu_init(int argc, char **argv, char **envp) {
    const size_t size = __init_array_end - __init_array_start;
    for (size_t i = 0; i < size; i++)
        (*__init_array_start[i])(argc, argv, envp);
}
```

# LIBC SPELUNKING

```
// Paraphrased from glibc/csu/elf-init.c
typedef void (*init_func)(int, char **, char **);

extern init_func __init_array_start[];
extern init_func __init_array_end[];

int __libc_csu_init(int argc, char **argv, char **envp) {
    const size_t size = __init_array_end - __init_array_start;
    for (size_t i = 0; i < size; i++)
        (*__init_array_start[i])(argc, argv, envp);
}
```

# LIBC SPELUNKING

```
// Paraphrased from glibc/csu/elf-init.c
typedef void (*init_func)(int, char **, char **);

extern init_func __init_array_start[];
extern init_func __init_array_end[];

int __libc_csu_init(int argc, char **argv, char **envp) {
    const size_t size = __init_array_end - __init_array_start;
    for (size_t i = 0; i < size; i++)
        (*__init_array_start[i])(argc, argv, envp);
}
```

OK BUT...

`__init_array_start?`  
`__init_array_end?`

OK BUT...

`__init_array_start?`  
`__init_array_end?`

[https://godbolt.org/z/cppcon2018\\_2](https://godbolt.org/z/cppcon2018_2)

# WHAT'S GOING ON HERE?

```
.section .init_array, "aw"
.align 8
.quad _GLOBAL__sub_I_Foo::numFoos
.text
```

# THE LINKER



# THE LINKER

What does it do?

- Resolves references between .o files
- Determines the layout of an executable
- Writes metadata

# A MORE REPRESENTATIVE PROGRAM

```
// hello.cpp
extern const char *getMessage();
void greet() {
    std::cout << getMessage() << "\n";
}
int main() {
    greet();
}
```

```
// message.cpp
const char *getMessage() {
    return "Hello world";
}
```

```
$ g++ -Os -o hello.o hello.cpp
$ g++ -Os -o message.o message.cpp
$ g++ -Os -o hello message.o hello.o
$ ./hello
Hello world
```

# OBJECT FILES

```
$ file hello hello.o message.o
hello: ELF 64-bit LSB executable, x86-64,
      dynamically linked, interpreter /lib64/ld-linux-x86-64.so.2,
      for GNU/Linux 3.2.0, not stripped
hello.o: ELF 64-bit LSB relocatable, x86-64, version 1 (SYSV), not stripped
message.o: ELF 64-bit LSB relocatable, x86-64, version 1 (SYSV), not stripped
```

# WHAT'S IN AN OBJECT FILE?

```
$ objdump -dC hello.o
```

# WHAT'S IN AN OBJECT FILE?

```
$ objdump -dC hello.o
```

```
research/out/hello/c++/none/dynamic/hello.o:      file format elf64-x86-64
```

```
Disassembly of section .text:
```

```
0000000000000000 <greet():>
 0: 55                      push   %rbp
 1: 48 89 e5                mov    %rsp,%rbp
 4: e8 00 00 00 00 00        callq  9 <greet()>+0x9>
 9: 48 89 c6                mov    %rax,%rsi
 c: 48 8b 05 00 00 00 00    mov    0x0(%rip),%rax          # 13 <greet()>+0x13>
13: 48 89 c7                mov    %rax,%rdi
16: e8 00 00 00 00          callq  1b <greet()>+0x1b>
1b: 48 8d 35 00 00 00 00    lea    0x0(%rip),%rsi          # 22 <greet()>+0x22>
22: 48 89 c7                mov    %rax,%rdi
25: e8 00 00 00 00          callq  2a <greet()>+0x2a>
```

# RELOCATIONS



# WHAT'S IN AN OBJECT FILE?

```
$ objdump --reloc -dC hello.o
```

```
research/out/hello/c++/none/dynamic/hello.o:      file format elf64-x86-64
```

```
Disassembly of section .text:
```

```
0000000000000000 <greet():
```

```
 0: 55                      push   %rbp
 1: 48 89 e5                mov    %rsp,%rbp
 4: e8 00 00 00 00           callq  9 <greet()>+0x9>
 5: R_X86_64_PLT32 getMessage()>-0x4
 9: 48 89 c6                mov    %rax,%rsi
 c: 48 8b 05 00 00 00 00    mov    0x0(%rip),%rax          # 13 <greet()>+0x13>
 f: R_X86_64_REX_GOTPCRELX std::cout->0x4
13: 48 89 c7                mov    %rax,%rdi
16: e8 00 00 00 00           callq  1b <greet()>+0x1b>
17: R_X86_64_PLT32 std::basic_ostream<char, std::char_traits<char> >& std::cout
```

# RELOCATIONS

- Different types
- Used within same object file

# SYMBOLS



# SYMBOLS

```
$ objdump --syms -C hello.o
```

# SYMBOLS

```
$ objdump --syms -C hello.o
```

```
research/out/hello/c++/none/dynamic/hello.o:      file format elf64-x86-64
```

## SYMBOL TABLE:

00000000	l	df	*ABS*	00000000	hello.cpp
00000000	l	d	.text	00000000	.text
00000000	l	d	.data	00000000	.data
00000000	l	d	.bss	00000000	.bss
00000000	l	d	.rodata	00000000	.rodata
00000000	l	o	.rodata	00000001	std::piecewise_construct
00000000	l	o	.bss	00000001	std::__ioinit
0000003d	l	F	.text	00000049	_static_initialization_and_destruction_0(int, i...
00000086	l	F	.text	00000015	_GLOBAL__sub_I_hello.cpp
00000000	l	d	.init_array	00000000	.init_array
00000000	l	d	.note.GNU-stack	00000000	.note.GNU-stack
00000000	l	d	.eh_frame	00000000	.eh_frame

# SYMBOLS

```
$ objdump --syms -C message.o
```

```
research/out/hello/c++/none/dynamic/message.o:      file format elf64-x86-64
```

## SYMBOL TABLE:

00000000	l	df	*ABS*	00000000	message.cpp
00000000	l	d	.text	00000000	.text
00000000	l	d	.data	00000000	.data
00000000	l	d	.bss	00000000	.bss
00000000	l	d	.rodata	00000000	.rodata
00000000	l	d	.note.GNU-stack	00000000	.note.GNU-stack
00000000	l	d	.eh_frame	00000000	.eh_frame
00000000	l	d	.comment	00000000	.comment
00000000	g	F	.text	00000008	getMessage()

# LINKER...

- Reads all the inputs
- Identifies symbols
- Applies relocations

*message.o*

.text  
getMessage()

.ro-data  
"Hello world"

*hello.o*

.text  
greet()

.text.init  
main()

*message.o*

.text  
getMessage()

.ro-data  
"Hello world"

*hello.o*

.text  
greet()

.text.init  
main()

*hello*

*Program Headers*

.text  
greet() {}  
getMessage()  
main()

.ro-data  
"Hello world"

# LINKER SCRIPTS

```
$ g++ -o /dev/null -x c /dev/null -Wl, --verbose
```

# LINKER SCRIPTS

```
$ g++ -o /dev/null -x c /dev/null -Wl,--verbose
```

```
GNU ld (GNU Binutils for Ubuntu) 2.30
```

```
Supported emulations:
```

```
elf_x86_64
```

```
elf32_x86_64
```

```
elf_i386
```

```
elf_iamcu
```

```
i386linux
```

```
elf_l1om
```

```
elf_k1om
```

```
i386pep
```

```
i386pe
```

```
using internal linker script:
```

```
=====
```

```
/* Script for -pie -z combreloc -z now -z relro: position independent executable,...  
/* Copyright (C) 2014-2018 Free Software Foundation, Inc.
```

```
.init_array      :  
{  
    PROVIDE_HIDDEN (__init_array_start = .);  
    KEEP (*SORT_BY_INIT_PRIORITY(.init_array.*)) SORT_BY_INIT_PRIORITY(.ctors.*));  
    KEEP (*(.init_array EXCLUDE_FILE (  
        *crtbegin.o *crtbegin?.o *crtend.o *crtend?.o ) .ctors))  
    PROVIDE_HIDDEN (__init_array_end = .);  
}
```

```
.init_array      :  
{  
    PROVIDE_HIDDEN (__init_array_start = .);  
    KEEP (*(SORT_BY_INIT_PRIORITY(.init_array.*)) SORT_BY_INIT_PRIORITY(.ctors.*)))  
    KEEP (*(.init_array EXCLUDE_FILE (  
            *crtbegin.o *crtbegin?.o *crtend.o *crtend?.o ) .ctors))  
    PROVIDE_HIDDEN (__init_array_end = .);  
}
```

```
.init_array      :  
{  
    PROVIDE_HIDDEN (__init_array_start = .);  
    KEEP (*(SORT_BY_INIT_PRIORITY(.init_array.*)) SORT_BY_INIT_PRIORITY(.ctors.*)))  
    KEEP (*(.init_array EXCLUDE_FILE (  
            *crtbegin.o *crtbegin?.o *crtend.o *crtend?.o ) .ctors))  
    PROVIDE_HIDDEN (__init_array_end = .);  
}
```

```
.init_array      :  
{  
    PROVIDE_HIDDEN (__init_array_start = .);  
    KEEP  (*($$SORT_BY_INIT_PRIORITY(.init_array.*)) SORT_BY_INIT_PRIORITY(.ctors.*)))  
    KEEP  (*(.init_array EXCLUDE_FILE (  
        *crtbegin.o *crtbegin?.o *crtend.o *crtend?.o ) .ctors))  
    PROVIDE_HIDDEN (__init_array_end = .);  
}
```

# NOW WE KNOW!

- Compiler:

# NOW WE KNOW!

- Compiler:
  - "static init" function for each TU
  - pointer to this function into `init_array`

# NOW WE KNOW!

- Compiler:
  - "static init" function for each TU
  - pointer to this function into `init_array`
- Linker:
  - gathers all `init_arrays` together
  - script defines symbols pointing at begin and end of `init_array`

# NOW WE KNOW!

- Compiler:
  - "static init" function for each TU
  - pointer to this function into `init_array`
- Linker:
  - gathers all `init_arrays` together
  - script defines symbols pointing at begin and end of `init_array`
- C runtime walks `init_array` and calls each

## STUFF TO KNOW

- You can write your own linker scripts
- Linker can discard unused sections: `-Wl,--gc-sections`
- Compiler flags: `-ffunction-sections`, `-fdata-sections`

# DYNAMIC LINKING

```
$ ls -l dynamic/hello static/hello
8,688      dynamic/hello*
2,406,632  static/hello*
```

# ANOTHER HELLO WORLD

```
$ g++ -Os -o message.o message.cpp
$ g++ -shared -o libhello.so message.o
$ g++ -Os -o hello hello.o
$ g++ -Os -o hello.o hello.cpp -L. -lhello
$ ./hello
Hello world
```

# MORE ELF HEADERS

```
$ readelf --dynamic --program-headers dynamic-dso/hello
```

```
Elf file type is EXEC (Executable file)
Entry point 0x4006f0
There are 9 program headers, starting at offset 64
```

## Program Headers:

Type	Offset	VirtAddr	PhysAddr	Flags	Align
	FileSiz	MemSiz			
PHDR	0x00000040	0x00400040	0x00400040		
	0x000001f8	0x000001f8	R E	0x8	
INTERP	0x00000238	0x00400238	0x00400238		
	0x0000001c	0x0000001c	R	0x1	
	[Requesting program interpreter: /lib64/ld-linux-x86-64.so.2]				
LOAD	0x00000000	0x00400000	0x00400000		
	0x00000ab0	0x00000ab0	R E	0x200000	
LOAD	0x00000da8	0x00600da8	0x00600da8		

# CODE ARCHAEOLOGY - PART 2



```
getMessage()@plt:  
0x4006b0: jmpq *0x200962(%rip) # 0x601018  
0x4006b6: pushq $0x0  
0x4006bb: jmpq 0x4006a0  
...  
0x601018: .quad 0x4006b6
```

```
getMessage()@plt:  
0x4006b0: jmpq *0x200962(%rip) # 0x601018  
0x4006b6: pushq $0x0  
0x4006bb: jmpq 0x4006a0  
...  
0x601018: .quad 0x4006b6
```

```
getMessage()@plt:  
0x4006b0: jmpq *0x200962(%rip) # 0x601018  
0x4006b6: pushq $0x0  
0x4006bb: jmpq 0x4006a0  
...  
0x601018: .quad 0x4006b6
```

```
getMessage()@plt:  
0x4006b0: jmpq *0x200962(%rip) # 0x601018  
0x4006b6: pushq $0x0  
0x4006bb: jmpq 0x4006a0  
...  
0x601018: .quad 0x4006b6
```

```
getMessage()@plt:  
0x4006b0: jmpq *0x200962(%rip) # 0x601018  
0x4006b6: pushq $0x0  
0x4006bb: jmpq 0x4006a0 ; ultimately resolves symbol 0  
...  
0x601018: .quad 0x4006b6
```

```
getMessage()@plt:  
0x4006b0: jmpq *0x200962(%rip) # 0x601018  
0x4006b6: pushq $0x0  
0x4006bb: jmpq 0x4006a0  
...  
0x601018: .quad 0x7ffff7bd35d5 ; now resolved to getMessage()
```

# DEBUGGING

- LD\_BIND\_NOW (and -Wl , -znow)
- ldd and LD\_DEBUG
- LD\_PRELOAD

# I WISH I HAD MORE TIME

- Weak references
- ODR violations
- LTO

## MORE READING

- Ian Lance Taylor's blog - [www.airs.com/blog](http://www.airs.com/blog)
- Honza Hubička's blog - [hubicka.blogspot.com/](http://hubicka.blogspot.com/)

# SPECIAL TRAINING EVENT

- Summer 2019 — Denver Area
- Charley Bay, Jason Turner and me together for 3 days
- C++20, error handling and performance
- Check out [coloradoplusplus.info](http://coloradoplusplus.info) for more info