

UEFI APPLICATIONS WITH C++

Morris Hafner (@mmhafner)

ABOUT ME

- Moved from Germany to Scotland
- Software engineer at Codeplay Software
- Postgrad student at University of Edinburgh
- Tendency to break compilers

WHAT IS (U)EFI?

- “New” firmware interface for x86, ARM, RISC-V
- Many features compared to BIOS
 - Network, Secure Boot, NVRAM, EFI Bytecode, ...
- No assembly required for many applications (e. g. bootloaders)

THE OLD WAYS: BIOS

- Loads 512B from MBR and jumps to 0x7C00
- Requires assembly code
- Bootloader switches to Protected Mode (32 Bit) and Long Mode (64 Bit)
- Bootloader scans ACPI, PCI, SMBus...
- May invoke BIOS services via interrupts

EFI EXECUTABLES

- COFF file format, subsystems 10-13
- MS calling convention
- UTF-16 strings
- On a FAT32 partition with partition type 0xEF00
- Implementations usually default to executing `EFI/Boot/bootx64.efi`

FIRST EXECUTABLE

TOOLCHAIN

- TianoCore EDK II, GNU efilib
- We will use GNU efilib, it is easier to use
- Available compilers: MSVC, MinGW gcc
- No clang
 - Doesn't support freestanding COFF executables

STANDARD LIBRARY

- no C or C++ standard library
 - A C95 stdlib implementation is part of EDK II
- Heap not easily accessible
- In practice, many headers can still be used
 - most of `<algorithm>`, `<array>`, `<tuple>`, `<type_traits>`

COMPILER INVOCATION

```
x86_64-w64-mingw32-g++ \  
-mno-red-zone \  
-ffreestanding -fshort-wchar \  
-nostdlib -e efi_main \  
-Wl,-dll -shared -Wl,--subsystem,10 \  
-c main.cpp
```

This can be wrapped in a CMake Toolchain

THE CODE

```
#include <efi.h>
#include <efilib.h>

extern "C" [[gnu::ms_abi]]
EFI_STATUS efi_main(
    IN EFI_HANDLE ImageHandle,
    IN EFI_SYSTEM_TABLE *SystemTable)
{
    SystemTable->ConOut->OutputString(
        SystemTable->ConOut,
        (CHAR16 *) L"Hello World\r\n"); // Note the missing const
}
```

RUNNING IT

OVMF is an OSS implementation that can be used with
QEMU

```
qemu-system-x86_64 \  
-drive file=hdd.img,if=ide \  
-bios OVMF.fd
```

ACCESSING EFI SERVICES

THE EFI PROTOCOL INTERFACE

- Firmware services can be queried using GUIDs
- Everything is loaded into structs
- OO style interface, but in C with function pointers

```
EFI_STATUS
LocateDevicePath (
    IN      EFI_GUID          *Protocol,
    IN OUT  EFI_DEVICE_PATH_PROTOCOL **DevicePath,
    OUT     EFI_HANDLE        *Device
);
```

HMMM...

Signature is usually

EFI_STATUS(in..., inout *..., out *...)

In C++, we want this:

```
expected<tuple<out...>, EFI_STATUS> const res =  
    func(in..., inout &...);  
if(res) {  
    auto[out...] = res.value();  
}
```

WHAT TO DO?

```
template <size_t in_count, size_t inout_count,  
         size_t out_count, typename Callable>  
auto wrap(Callable c);  
  
auto locate_device_path = wrap<1, 1, 1>(LocateDevicePath);  
expected<tuple<EFI_HANDLE>, EFI_STATUS> result =  
    locate_device_path(protocol, device_path);
```

1. Partition the argument list
2. Create a new function wrapping the EFI function
3. Wrap the error codes

GETTING THE ARGUMENT LIST

→ Boost Callable Traits

```
EFI_STATUS
LocateDevicePath(EFI_GUID *,
    EFI_DEVICE_PATH_PROTOCOL **,
    EFI_HANDLE *);

#include <boost/callable_traits/args.hpp>
using args =
    boost::callable_traits::args_t<decltype(LocateDevicePath)>;
static_assert(is_same_v<args,
    tuple<EFI_GUID *,
        EFI_DEVICE_PATH_PROTOCOL **,
        EFI_HANDLE *>
    >);
```


PARTITION THE ARGUMENT LIST

```
template <size_t offset, typename T, size_t... I>
auto constexpr st_impl(T tuple, index_sequence<I...>) {
    return tuple{get<offset + I>(tuple)...};
}
```

```
template <size_t N, typename T>
auto constexpr split_tuple(T tuple) {
    auto constexpr sz = tuple_size_v<T>;
    return tuple{
        st_impl<0>(tuple, make_index_sequence<N>{}),
        st_impl<N>(tuple, make_index_sequence<sz - N>{}),
    };
}
```

3-WAY SPLIT

```
using in_split =  
    decltype(split_tuple<in_count>(declval<args>()));  
using In = // <- first  
    tuple_element_t<0, in_split>;  
using InOut_Out =  
    tuple_element_t<1, in_split>;  
using inout_out_split =  
    decltype(split_tuple<inout_count>(declval<InOut_Out>()));  
using InOut = // <- second  
    tuple_element_t<0, inout_out_split>;  
using Out = // <- third  
    tuple_element_t<1, inout_out_split>;
```

WRAP THE FUNCTION

```
template <typename Func,  
    typename... Is, typename... IOs, typename... Os>  
auto constexpr make_out_param_adapter(  
    Func func, tuple<Is...>, tuple<IOs...>, tuple<Os...>) {  
    return [func](Is... is, IOs... ios) {  
        tuple<remove_pointer_t<Os>...> res;  
        auto const ptr = transform_tuple(res, [] (auto &val) {  
            return addressof(val);  
        });  
        apply(c, tuple_cat(  
            tuple{get_ptr(is)...}, tuple{get_ptr(ios)...}, ptr));  
        return res;  
    });  
}
```

OVERHEAD?

OVERHEAD!

`std::tuple` causes value initialization

```
template <typename T> struct uninitialized {  
    uninitialized() {}  
    T & get() {  
        return val;  
    }  
    T val;  
};
```

```
tuple<uninitialized<remove_pointer_t<Os>>...> res;  
auto const ptr = transform_tuple(res, [] (auto &val) {  
    return val.get();  
});
```

WRAP THE ERROR CODES

- `expected<T, E>` contains either a value or an error
- We use Simon Brands `tl::expected<T, E>`

```
if constexpr(
    is_same_v<invoke_result_t<Func, Is..., IOs..., Os...>,
              EFI_STATUS>) {
    using result_t =
        tl::expected<tuple<remove_pointer_t<Os>...>, EFI_STATUS>;
    if(auto result = apply(...) != EFI_SUCCESS)
        return result_t(tl::unexpected, result);
    return result_t{transform_tuple(...)};
} else { // void
    apply(...);
    return res;
}
```

OVERHEAD?

WE MADE SIMPLIFICATIONS

- Assumptions that there are only fundamental types and PODs
 - Types must be trivially constructible (`uninitialized<T>`)
 - Compiler can go much further thanks to trivial copy construction and destruction
- No overloads (Callable Traits)
 - Not too bad, C doesn't have overloads or destructors
- But may be wrapped by a future `std::overload`

EXAMPLE APPLICATION

RENDERING TO THE FRAMEBUFFER

- We could write our own kernel in C++ now...
 - `ExitBootServices()`
- Let's render a couple of spheres onto the framebuffer instead

CREATING A GRAPHICS OUT PROTOCOL INSTANCE

First, some wrappers

```
// Look for GOP implementations  
auto locate_handle_buffer =  
    wrap<3, 1, 1>(bootServices->LocateHandleBuffer);  
// Create a GOP Instance  
auto handle_protocol =  
    wrap<2, 0, 1>(bootServices->HandleProtocol);
```

YAY, MONADS!

```
uint64_t handle_count = 0;
auto maybe_gop =
locate_handle_buffer(ByProtocol, graphicsOutProtocolGUID,
                     nullptr, handle_count)
    .and_then([&](auto proto_impl) {
        return handle_protocol(get<0>(proto_impl)[0],
                                graphicsOutProtocolGUID);
    })
    .map([](auto opened_gop) {
        return reinterpret_cast<EFI_GRAPHICS_OUTPUT_PROTOCOL*>(
            get<0>(opened_gop));
    });
if(!maybe_gop)
    conOut->OutputString(conOut, (CHAR16*) u"Fail\r\n");
```

FACT: This slide contains error handling

CREATE THE FRAMEBUFFER

```
gop = maybe_gop.value();
auto query_mode = wrap<2, 0, 2>(gop->QueryMode);
for(int i = 0;; ++i) {
    auto mode_info = query_mode(gop, i)
    .map([](auto r)->decltype(auto){return *std::get<1>(r);});
    if(!mode_info) break;
    if(mode_info->HorizontalResolution == width &&
        mode_info->VerticalResolution == height &&
        mode_info->PixelFormat ==
        PixelBlueGreenRedReserved8BitPerColor) {
        gop->SetMode(gop, i);
        break;
    }
}
```

USE THE FRAMEBUFFER

With emulated double buffering

```
void swap_to_screen() const {
    auto const pixel_ptr =
        reinterpret_cast<uint32_t*>(gop->Mode->FrameBufferBase);
    for(auto const& row : rows) {
        std::copy(row.begin(), row.end(), pixel_ptr);
        pixel_ptr += gop->Mode->Info->PixelsPerScanLine;
    }
}

void clear() {
    for(auto &row : rows) {
        std::fill(row.begin(), row.end(), T{});
    }
}
```

IMPLICIT SURFACES

Remember: Avoided heap

Implicit surfaces using signed distance fields offer a functional representation of a scene

Unfortunately no time for details

```
auto scene = [](vec3 const& rayT) {  
    return pUnion(sphere(vec3( 10.0, 0.0, 50), 20, rayT),  
                  sphere(vec3(-10.0, 5.0, 50), 20, rayT));  
};
```

CORRECT CODE?

We are in a freestanding environment, yet we use features that are not available there.

See Ben Craigs P0829 “Freestanding Proposal” for a solution.

**A FUTURE WITH ZERO
COST EXCEPTIONS**

HERB SUTTER: P0709

“Zero-overhead deterministic exceptions”

```
constexpr auto make_out_param_adapter() {  
    return [] (args...) -> std::tuple<outs...> throws(EFI_STATUS) {  
        // ...  
        if(result != 0) {  
            throw result;  
        }  
        // ...  
    }  
}
```

REFERENCES

- github.com/mmha/efiraytracer
- github.com/TartanLlama/expected
- OSDev Wiki: UEFI Bare Bones
- P0829r2: Freestanding Proposal